

# VeriTrace: Towards Automatic Testing and Verification of Concurrent Programs

ZHANG Yu

LCS, ISCAS



*Suzhou • November 2012*

# Motivation

- Concurrent programs are notoriously hard to write correctly.
- Testing and debugging of concurrent programs are also hard.
  - Bugs may not recur due to non-determinism.
  - Enumerating all possible executions is hard or practically impossible.
- There are not many practical tools for testing and verifying concurrent programs.
- We also need a tool for teaching purpose at CAS.

# Concurrent Objects

- Our tool runs on **JVM**.
- We talk about **objective** concurrent programs
  - Concurrent objects are shared by multiple processes or threads.
  - Threads can call methods to read information and make change to the shared object.
  - Method execution takes time, and execution time in different threads can overlap.
- What does it mean for concurrent objects and methods to be **correct**?

# Correctness for Concurrent Programs


- Sequential consistency
  - Every concurrent execution has a consistent sequential execution.
  - The sequential execution preserves the program/execution order in every single process.
- Linearizability
  - Stricter than sequential consistency
  - The sequential execution preserves the happen-before order between **all process**.
- Quiescent consistency
- Looser consistency at hardware/architecture level.

# Linearizability

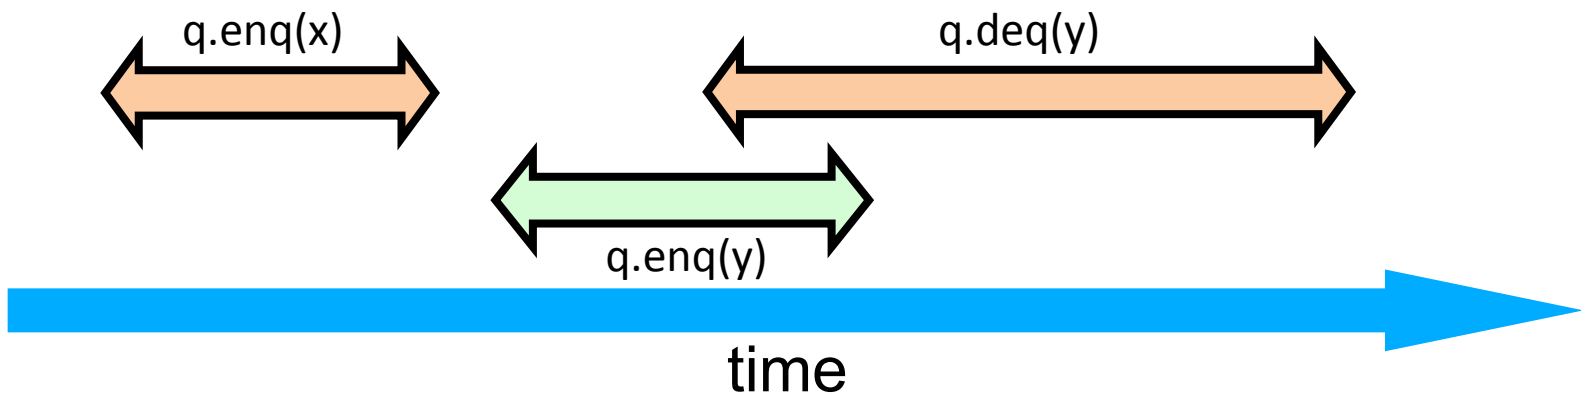
- Linearizable execution
  - Every method should “take effect” **instantaneously** between invocation and response.
  - Concurrent execution is correct if this sequential execution is correct.
- A concurrent object is linearizable if all possible executions are linearizable.

# Linearizability

- Linearizable execution:




- Non-linearizable execution:



# Trace Model

- We record traces of method executions in all threads
  - Every method execution has two events
    - Method call/invocation: method name & arguments
    - Method return/response: result or exception
  - A **concurrent trace** records method events tagged by thread ID, in temporal order.
  - Happen-before: a method execution  $m_1$  **happens before**  $m_2$  if  $m_1$ 's response is before  $m_2$ 's invocation in the trace.

B:q.enq(y)  
A:q.enq(x)  
B:q.enq:void  
B:q.deq()  
A:q.enq:void  
B:q.deq:x  
A:q.deq()  
A:q.deq:y

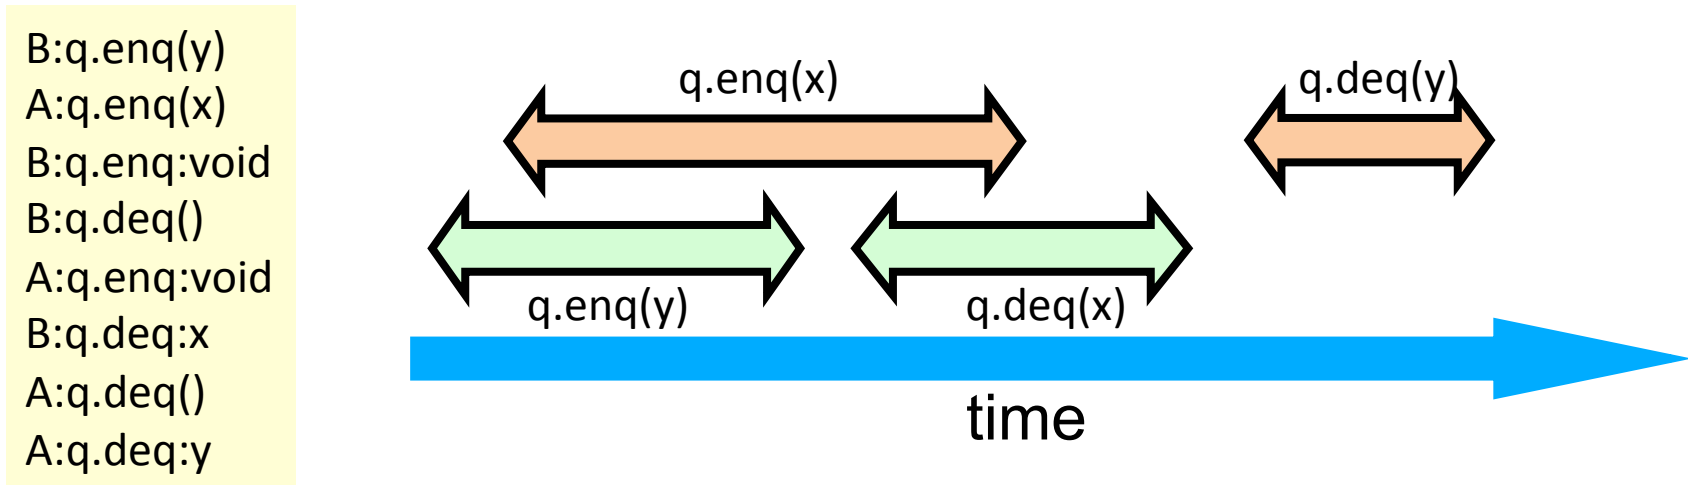


# Serialization

- A **simulation trace** is a sequence of method calls and arguments.
  - E.g., {q.enq(x), q.enq(y), q.deq(), q.deq()}
  - Simulation traces are intended to be executed in a single-process mode.
- Given a concurrent trace, **serialization** produces all its possible simulation traces, **preserving happen-before relation** between method executions.
  - Executing a simulation trace produces a **sequential trace**, which is a sequence of method calls **and results**.
  - A sequential trace **matches** its original concurrent trace if every method execution returns the same result as in the concurrent one.



# Serialization



- Its sequential traces are

A:q.enq(x):void  
B:q.enq(y):void  
B:q.deq(): x  
A:q.deq(): y

B:q.enq(y):void  
A:q.enq(x):void  
B:q.deq(): y  
A:q.deq(): x

B:q.enq(y):void  
B:q.deq(): y  
A:q.enq(x):void  
A:q.deq(): x

- Only the first trace is a matching trace.

# Correctness

- Correctness in our model:
  - A concurrent trace is correct if **its serialization has a matching sequential trace**.
  - A concurrent program is correct if all its concurrent traces are correct.

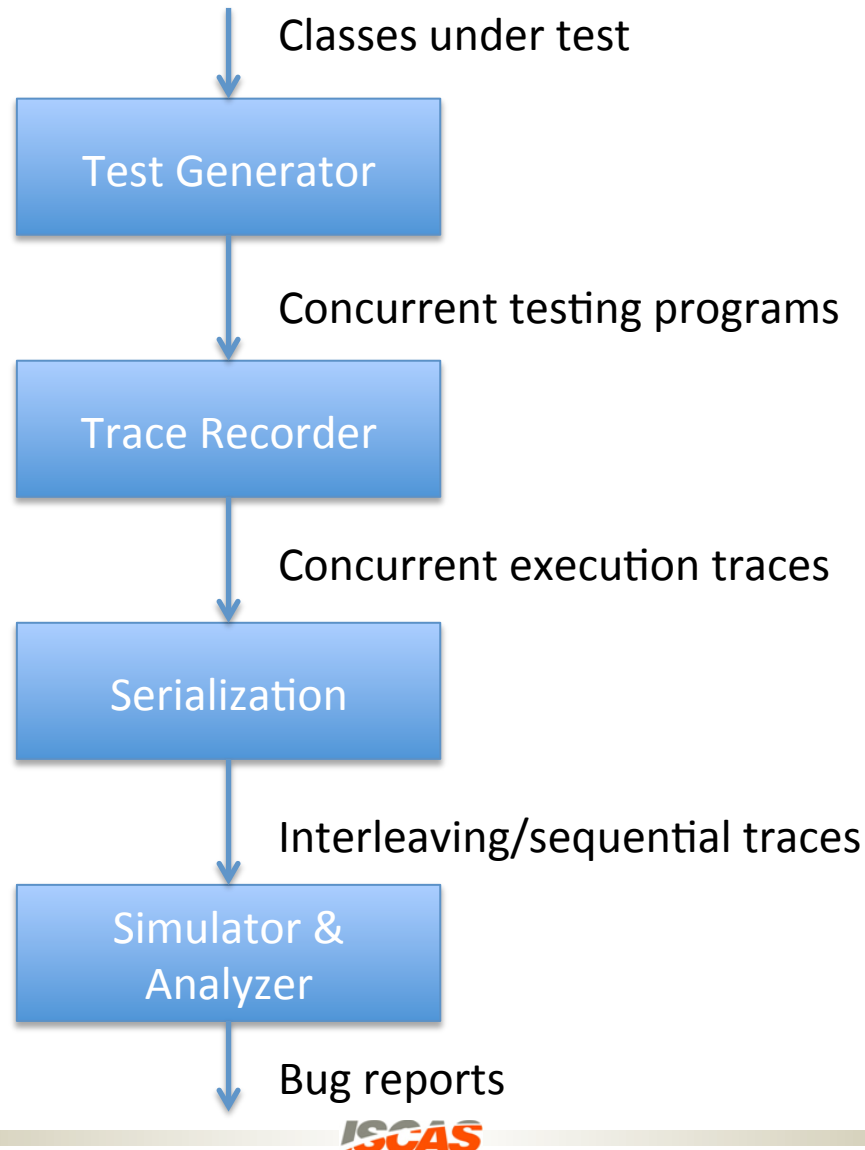
# Concurrent Trace Generation

- Enumerating all possible concurrent traces is hard or practically impossible.
- We generate traces via **testing**: given a concurrent class for verification, a test case includes
  - a number of threads, and
  - a random sequence of method calls (and arguments) for each thread.
- We can record concurrent traces for every test case at the JVM level.
  - **No annotation is required** for source programs.
  - Test and trace generation are **fully automatic**.

# Simulation and Analysis

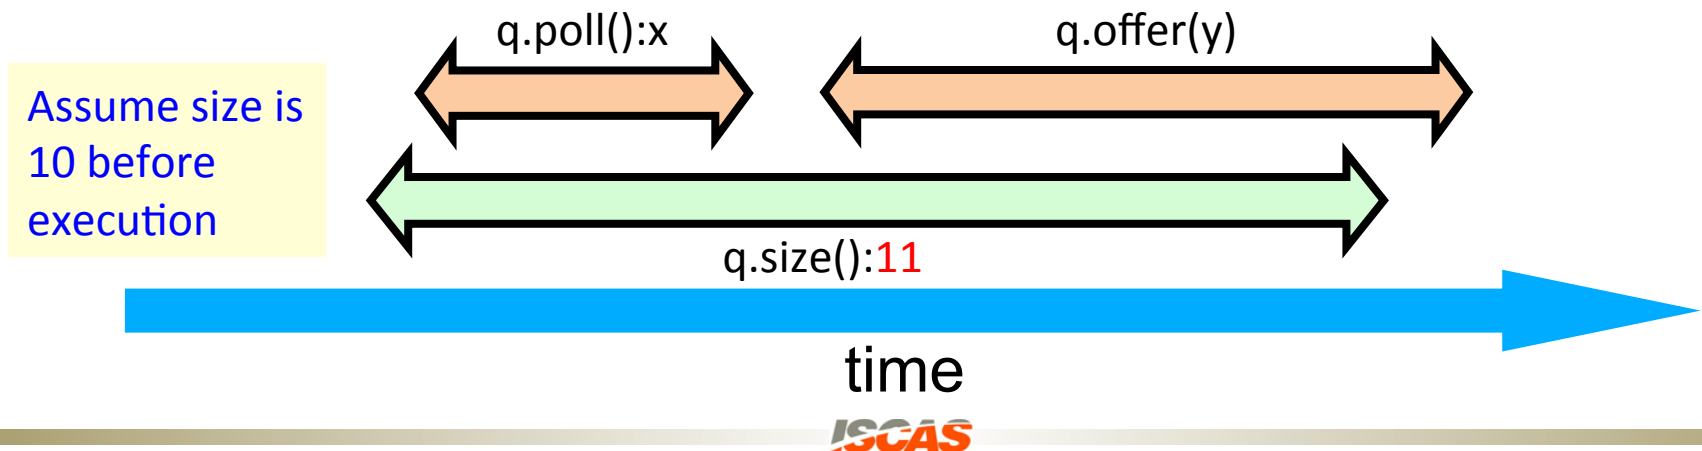
- With each concurrent trace,
  - Serialization: produce all its simulation traces.
  - Simulate: execute the methods in the simulation trace in a **single-thread** manner, and check the result against the concurrent trace.
  - If no matching sequential trace is found, report the test with the buggy trace.
- The report is **true negative**: we do not produce false alarm.

# Design of VeriTrace



# Experiments

- `java.util.concurrent.ConcurrentLinkedQueue` (2 threads)
  - Methods: `offer/poll/size`, 1000 tests
    - 200 method calls per thread, 2~3 buggy traces
    - 500 method calls per thread, 10~20 buggy traces
  - All buggy traces has the pattern: **size || poll; offer**
  - Method `size` is NOT thread-safe: it is stated in JSE7, but not in JSE6.




# Experiments

- LockFreeList (Chapter 9, “*The Art of Multiprocessor Programming*”)
  - Methods: add/remove, 1000 tests
    - 200 method calls per thread,  $\leq 10$  buggy traces
    - 500 method calls per thread, 30 ~ 40 buggy traces
  - All buggy traces has the pattern: **remove || remove**
  - Method remove is buggy, which is indicated by the online errata of the book.
  - The report also shows that data race only occurs between removes.

# Implementation Issues

- Shallow simulation vs. deep simulation
  - A long concurrent trace often presents the following pattern:



- **Deep simulation** produces full traces and starts simulation from scratch.
  - With a lazy interleaving algorithm it does not require much memory.
- **Shallow simulation** do simulation for each segment and records intermediate states.
  - It often finds bugs quickly, faster than deep simulation.



# Ongoing and Future Work

- The tool is still under intensive development.
- On-going work:
  - Use the tool to verify practical concurrent libraries.
  - Extend testing module to allow user-defined test cases.
  - Compare with other tools, e.g., LineUp, Thread-Safe.
- Future work:
  - Further analysis on buggy traces. Can help identify buggy methods?
  - Refined trace recoding to provide more information with buggy traces.
  - Combined with other program analysis techniques, e.g., model-checking, static analysis, etc.